

如何快速上手 Kubernetes

张晋涛

个人介绍

- 张晋涛 @API7.AI
- Apache APISIX Committer
- Kubernetes Ingress-nginx reviewer
- containerd/Docker/Helm/Kubernetes/KIND contributor
- 『K8S 生态周报』维护者
- 《Kubernetes上手实践》小册作者
- <https://github.com/tao12345666333>



Agenda

- Kubernetes 概览
- Kubernetes 环境准备
- Kubernetes 中的应用部署
- Kubernetes 中的流量路由



Kubernetes 概览

Kubernetes 发展历程

- 2014 年由 Google 开源
- 大规模场景下 Docker 容器编排
- 隔离性 / 标准化
- 滚动更新
- 故障自愈
- 动态扩 / 缩容



Kubernetes 发展历程

- CNCF 首个毕业项目
- Kubernetes 在生产环境中的应用比例达到 83%
- GitHub star 数 80.9k

注：以上统计数据来自 CNCF 的 2020 调查报告



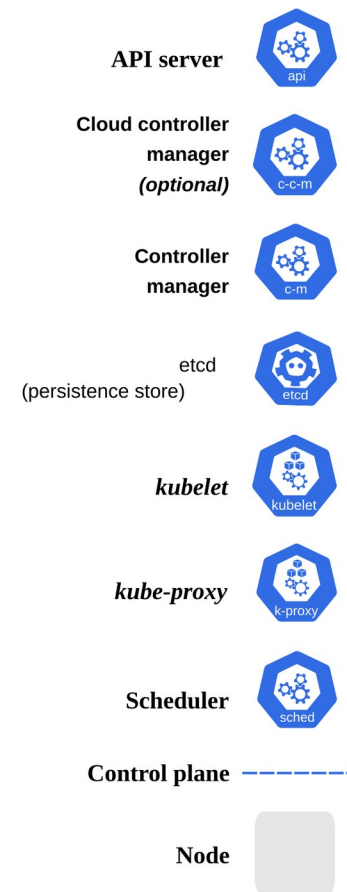
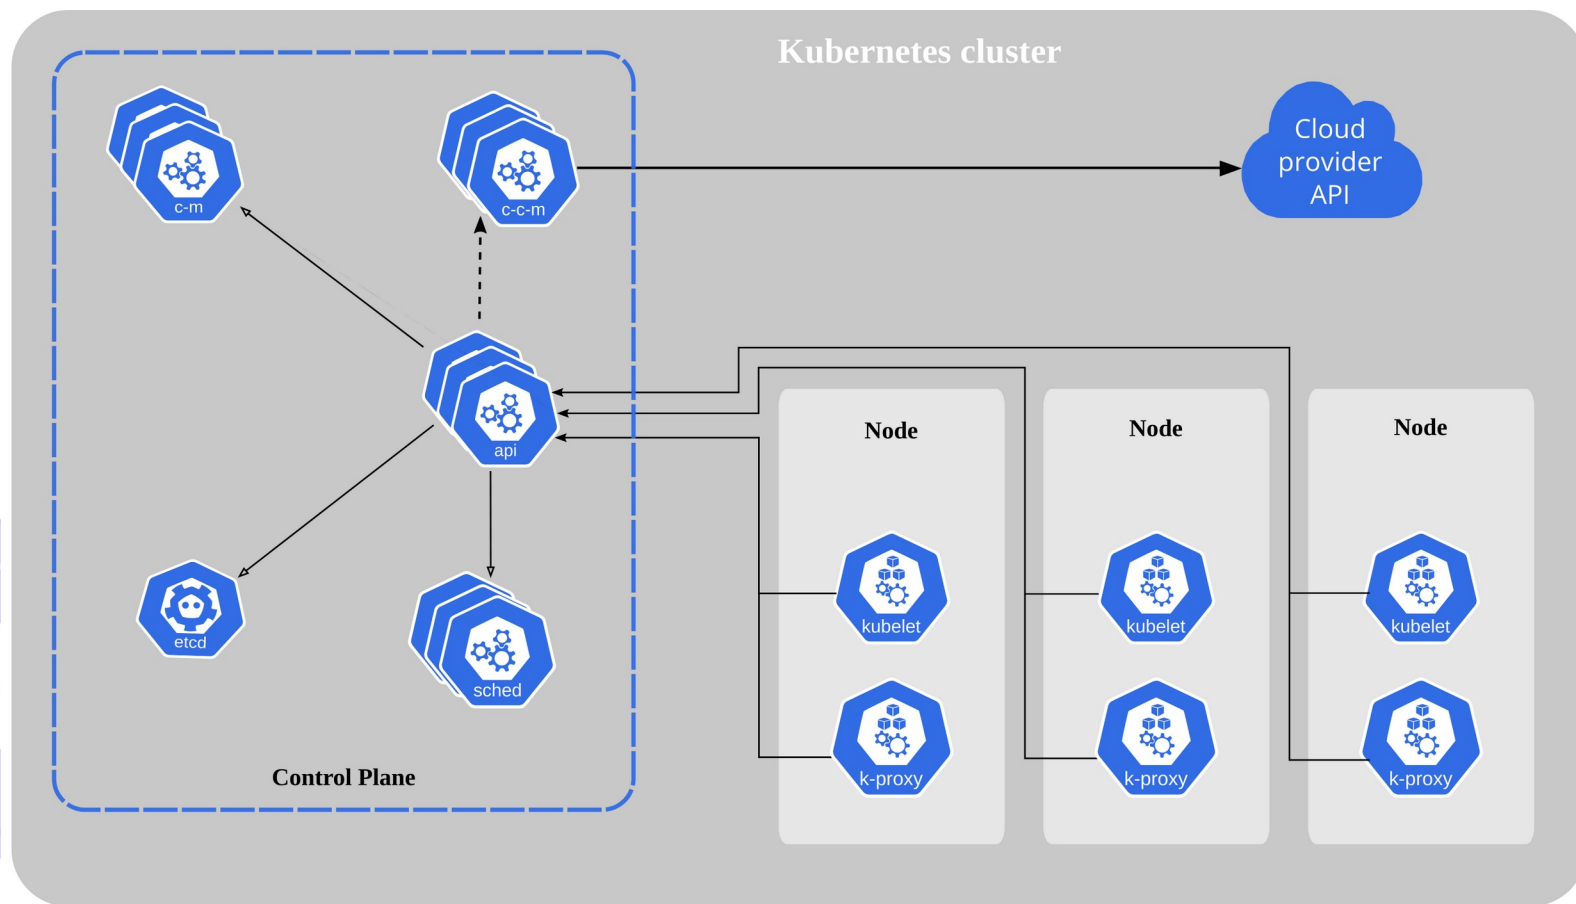
掘金小册 juejin.cn/books

为什么需要 Kubernetes

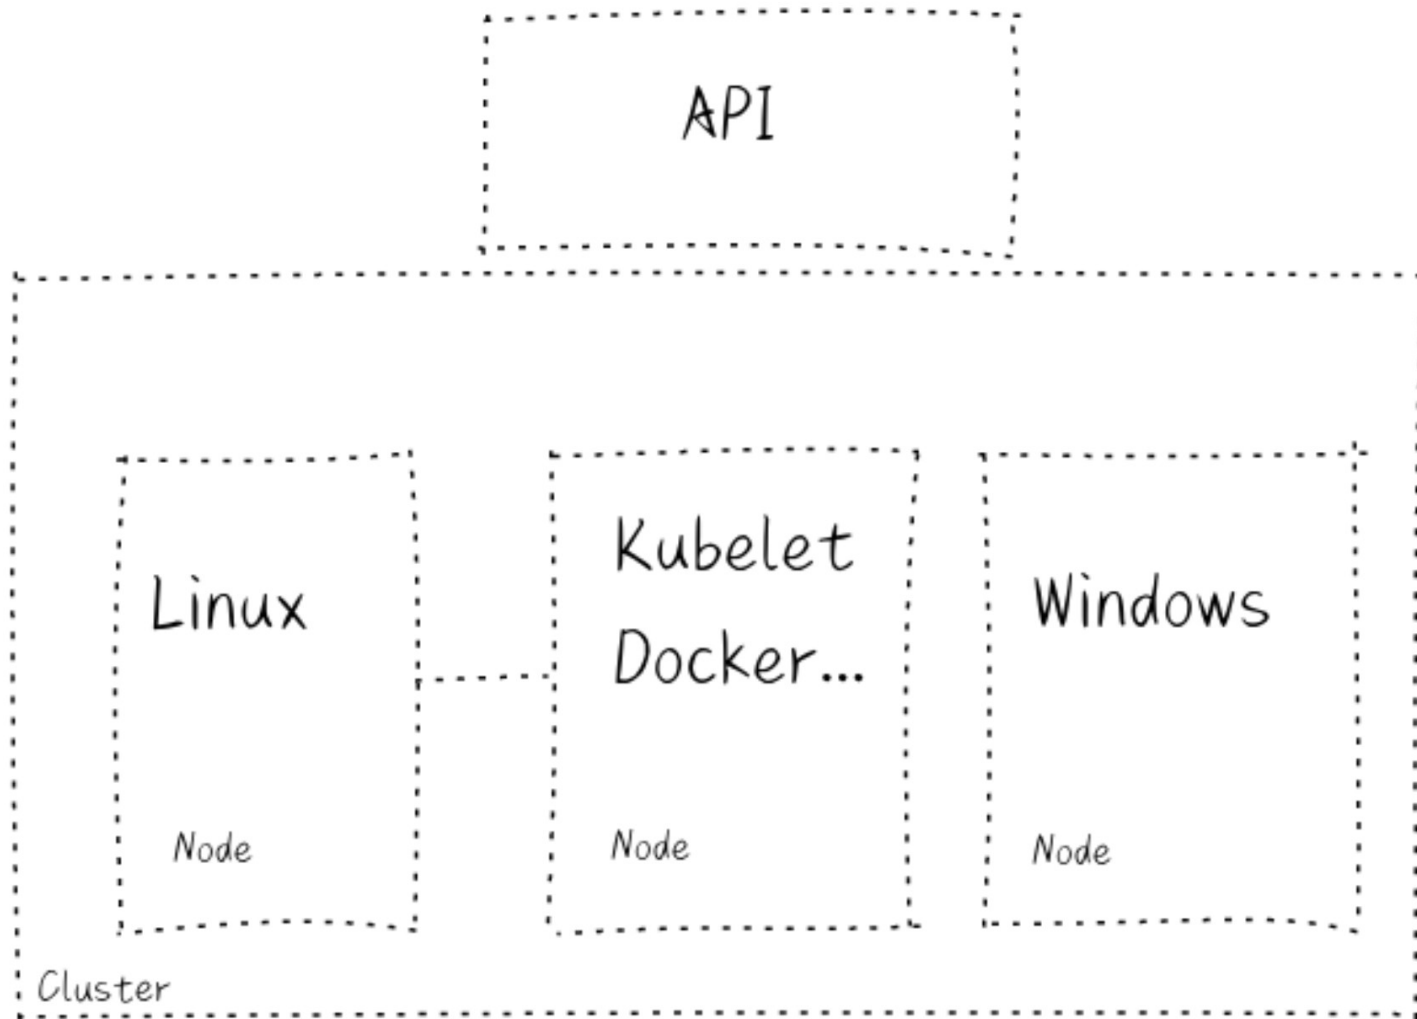
- 大规模多节点下的容器编排
- 快速扩 / 缩容
- 故障自愈
- 弹性
- 技术趋势
- 一致性
- 不锁定



Kubernetes 架构



理解 Kubernetes



Kubernetes 环境准备

在线环境

- Katacoda
- Play with Kubernetes



本地测试环境

- Docker Desktop （内置）
- KIND （ Kubernetes In Docker ）
- [Lima](#)
- Minikube
- Microk8s
- ...



生产部署

- kubectl (推荐)
- kops (AWS)
- Kubespray
- ...



使用 KIND 创建本地环境

- 准备 Docker 环境
- 下载 KIND 二进制文件



```
1 → ~ kind create cluster --name moelove.info
2 Creating cluster "moelove.info" ...
3 ✓ Ensuring node image (kindest/node:v1.21.1)
4 ✓ Preparing nodes
5 ✓ Writing configuration
6 ✓ Starting control-plane
7 ✓ Installing CNI
8 ✓ Installing StorageClass
9 Set kubectl context to "kind-moelove.info"
10 You can now use your cluster with:
11
12 kubectl cluster-info --context kind-moelove.info
13
14 Have a question, bug, or feature request? Let us know! https://kind.sigs.k8s.io/#community 😊
15 → ~ kubectl get nodes
16 NAME                                STATUS    ROLES                  AGE     VERSION
17 moelove.info-control-plane          Ready    control-plane,master   42s     v1.21.1
```

Kubernetes 中的应用部署

Kubernetes 中的核心概念

- Resource Object
 - Spec
 - Status
- Pod
- Deployment
- Daemonset
- Service
- Statefulset
- Job/CronJob
- Ingress
- Node



Kubernetes 核心设计理念

- 声明式
- 无侵入性
- 可移植（提供抽象）
- 显式定义

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: nginx-deployment
5 spec:
6   selector:
7     matchLabels:
8       app: nginx
9   minReadySeconds: 5
10  template:
11    metadata:
12      labels:
13        app: nginx
14    spec:
15      containers:
16        - name: nginx
17          image: nginx:1.14.2
18          ports:
19            - containerPort: 80
```



Kubernetes 如何管理容器

- Kubernetes 中最小调度单元是 Pod
- 每个 Pod 可以包含多个 container
- 每个 Pod 有自己的 IP
- 可通过集群网络与其他 Pod 通信
- Pod 内容器共享网络堆栈



Kubernetes 运行 Pod

```
1 → ~ kubectl run moelove-redis --image="redis:alpine" --restart=Never
2 pod/moelove-redis created
3 → ~ kubectl get pods
4 NAME          READY   STATUS    RESTARTS   AGE
5 moelove-redis  1/1     Running   0           31s
6 → ~ kubectl get pods -owide
7 NAME          READY   STATUS    RESTARTS   AGE   IP            NODE                                NOMINATED NODE
8 moelove-redis  1/1     Running   0           53s   10.244.0.5    moelove.info-control-plane        <none>
9 → ~ kubectl describe pods moelove-redis
10 Name:          moelove-redis
11 Namespace:     default
12 Priority:       0
13 Node:          moelove.info-control-plane/172.18.0.8
14 Start Time:    Tue, 14 Sep 2021 16:49:27 +0800
15 Labels:        run=moelove-redis
16 Annotations:   <none>
17 Status:        Running
18 IP:            10.244.0.5
19 IPs:
20 IP: 10.244.0.5
21 Containers:
22   moelove-redis:
23     Container ID: containerd://c2874f6a90cbf6f43224cda5f8677e8c03c96baf3105da17bad48c2b470eaef4
24     Image:        redis:alpine
25     Image ID:
26     docker.io/library/redis@sha256:fa785f9bd167b94a6b30210ae32422469f4b0f805f4df12733c2f177f500d1ba
27     Port:         <none>
28     Host Port:    <none>
29     State:        Running
30       Started:    Tue, 14 Sep 2021 16:49:47 +0800
31     Ready:        True
32     Restart Count: 0
```

Kubernetes 如何启动 Pod

```
1 → ~ kubectl get events
2 LAST SEEN    TYPE      REASON              OBJECT               MESSAGE
3 9m25s        Normal    Scheduled            pod/moelove-redis    Successfully assigned
    default/moelove-redis to moelove.info-control-plane
4 9m24s        Normal    Pulling             pod/moelove-redis    Pulling image "redis:alpine"
5 9m6s         Normal    Pulled              pod/moelove-redis    Successfully pulled image
    "redis:alpine" in 18.76340441s
6 9m6s         Normal    Created             pod/moelove-redis    Created container moelove-
    redis
7 9m5s         Normal    Started             pod/moelove-redis    Started container moelove-
    redis
```



Kubernetes 使用 Deployment 部署

```
1 → ~ kubectl create deploy moelove-redis1 --image="redis:alpine"
2 deployment.apps/moelove-redis1 created
3 → ~ kubectl get pods -o wide
4 NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE
  NOMINATED NODE   READINESS GATES
5 moelove-redis1-7ccc854d6f-hbxj2      1/1     Running   0           14s   10.244.0.6      moelove.info-control-plane
  <none>             <none>
6 → ~ kubectl get deploy
7 NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
8 moelove-redis1                      1/1     1             1           14s
```



Kubernetes 重启 Pod

```
1 → ~ kubectl get pods -o wide
2 NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE
  NOMINATED NODE   READINESS GATES
3 moelove-redis1-7ccc854d6f-hbxj2     1/1     Running   0           14s   10.244.0.6    moelove.info-control-plane
  <none>            <none>
4 → ~ kubectl rollout restart deploy/moelove-redis1
5 deployment.apps/moelove-redis1 restarted
6 → ~ kubectl get pods -owide
7 NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE
  NOMINATED NODE   READINESS GATES
8 moelove-redis1-7ff6f678bb-dbbqx     1/1     Running   0           4s    10.244.0.7    moelove.info-control-plane
  <none>            <none>
```



Kubernetes 使用 YAML 部署

```
1 → ~ kubectl create deploy moelove-redis2 --image="redis:alpine" --dry-run=client -o yaml
2 apiVersion: apps/v1
3 kind: Deployment
4 metadata:
5   creationTimestamp: null
6   labels:
7     app: moelove-redis2
8   name: moelove-redis2
9 spec:
10  replicas: 1
11  selector:
12    matchLabels:
13      app: moelove-redis2
14  strategy: {}
15  template:
16    metadata:
17      creationTimestamp: null
18      labels:
19        app: moelove-redis2
20    spec:
21      containers:
22        - image: redis:alpine
23          name: redis
24          resources: {}
25  status: {}
```



Kubernetes 中的流量路由

Kubernetes 中的 Service

```
1 → ~ kubectl create deploy httpbin1 --image='kennethreitz/httpbin'
2 deployment.apps/httpbin1 created
3 → ~ kubectl create deploy httpbin2 --image='kennethreitz/httpbin'
4 deployment.apps/httpbin2 created
5 → ~ kubectl expose deploy httpbin1 --port=80
6 service/httpbin1 exposed
7 → ~ kubectl expose deploy httpbin2 --port=80
8 service/httpbin2 exposed
9 → ~ kubectl get pod,svc -owide
```

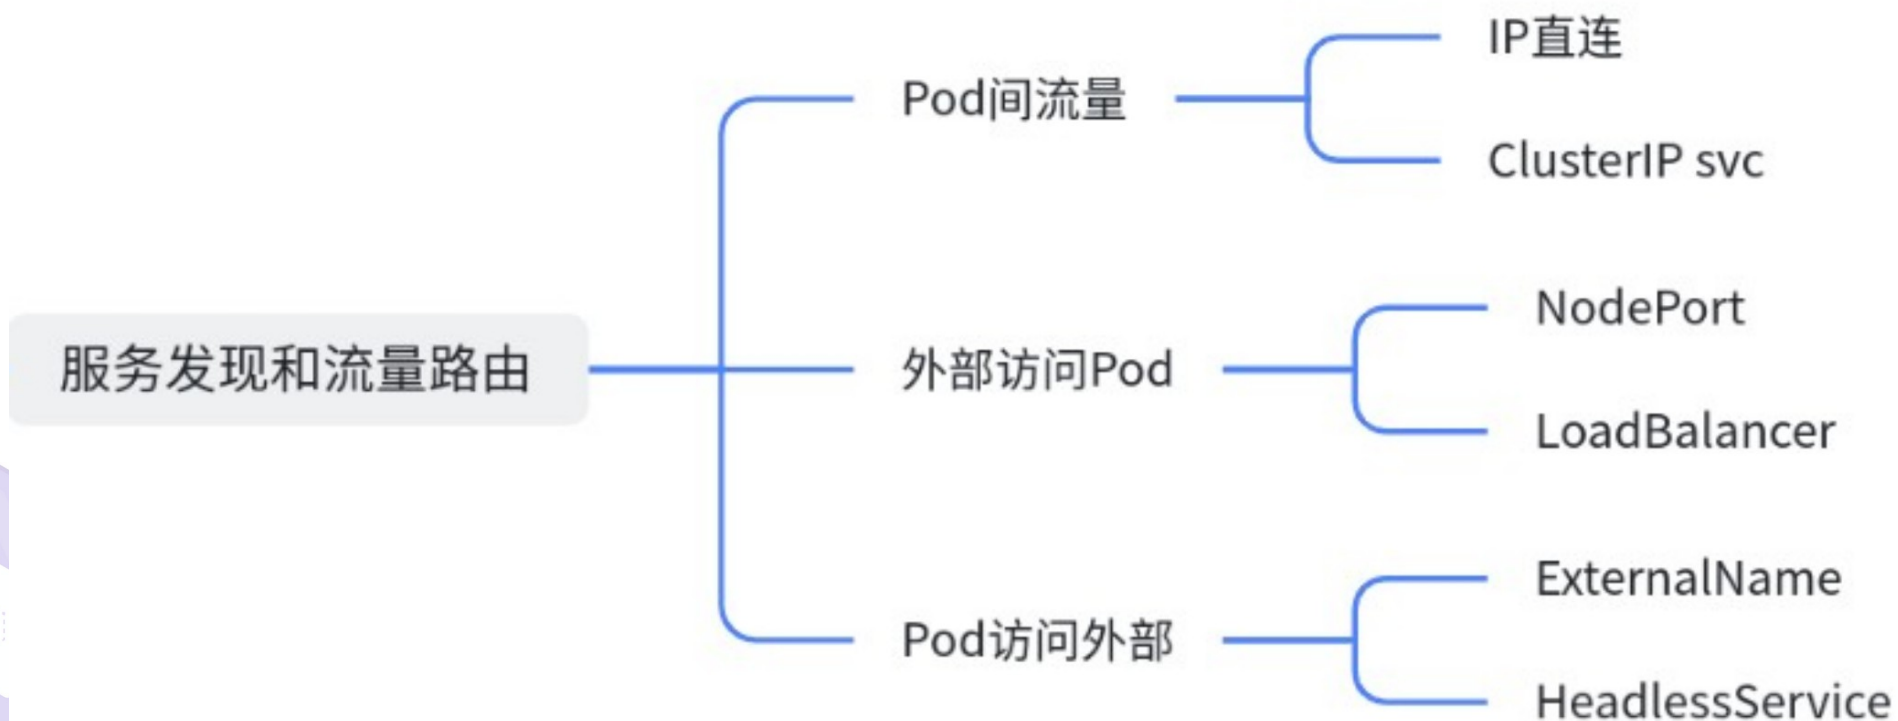
10	NAME		READY	STATUS	RESTARTS	AGE	IP	NODE
11	pod/httpbin1-56f8c5cf48-hcztx	1/1	Running	0	3m43s	10.244.0.8	moelove.info-control-plane	
	<none>	<none>						
12	pod/httpbin2-55bcc4cb8f-wg7tn	1/1	Running	0	3m39s	10.244.0.9	moelove.info-control-plane	
	<none>	<none>						
13								
14	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR	
15	service/httpbin1	ClusterIP	10.96.184.144	<none>	80/TCP	3m12s	app=httpbin1	
16	service/httpbin2	ClusterIP	10.96.113.237	<none>	80/TCP	3m8s	app=httpbin2	
17	service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	34m	<none>	

Kubernetes 中的 Service

```
1 → ~ kubectl exec -it httpbin1-56f8c5cf48-hcztx -- bash
2 root@httpbin1-56f8c5cf48-hcztx:/# curl httpbin2/ip
3 {
4   "origin": "10.244.0.8"
5 }
6
7 root@httpbin1-56f8c5cf48-hcztx:/# curl 10.244.0.9/ip
8 {
9   "origin": "10.244.0.8"
10 }
11
```



Kubernetes 中的 Service



使用 NodePort 在集群外访问服务

```
1 → ~ kubectl expose deploy httpbin2 --port=80 --type=NodePort --name httpbin2-nodeport
2 service/httpbin2-nodeport exposed
3 → ~ kubectl get svc
4 NAME                TYPE          CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
5 httpbin1             ClusterIP     10.96.184.144   <none>           80/TCP           19m
6 httpbin2             ClusterIP     10.96.113.237   <none>           80/TCP           19m
7 httpbin2-nodeport    NodePort      10.96.44.67     <none>           80:31145/TCP     5s
8 kubernetes           ClusterIP     10.96.0.1       <none>           443/TCP          50m
9 → ~ curl `kubectl get nodes -o jsonpath="{.items[0].status.addresses[0].address}"`:31145/ip
10 {
11   "origin": "10.244.0.1"
12 }
```



使用 Gateway 在集群外访问服务

```
1 → ~ kubectl create ns apisix
2 namespace/apisix created
3 → ~ helm install apisix apisix/apisix --set gateway.type=NodePort --set ingress-controller.enabled=true --
  namespace apisix
4 ...
5 → ~ kubectl -n apisix exec deploy/apisix -- curl -s http://httpbin2.default:80/ip
6 {
7   "origin": "10.244.0.10"
8 }
9 → ~ kubectl -n apisix exec deploy/apisix -- curl -s "http://127.0.0.1:9180/apisix/admin/routes/1" -H "X-API-
  KEY: edd1c9f034335f136f87ad84b625c8f1" -X PUT -d '
10 {
11   "uri": "/ip",
12   "host": "httpbin.local",
13   "upstream": {
14     "type": "roundrobin",
15     "nodes": {
16       "httpbin2.default:80": 1
17     }
18   }
19 }'
20 {"node":{"key":"\\apisix\\routes\\1","value":{"upstream":
  {"type":"roundrobin","scheme":"http","hash_on":"vars","nodes":
  {"httpbin2.default:80":1},"pass_host":"pass"},"update_time":1631613271,"host":"httpbin.local","priority":0,"cr
  eate_time":1631613217,"id":"1","uri":"\\ip","status":1}},"action":"set"}
21
```

使用 Gateway 在集群外访问服务

```
1 → ~ kubectl -n apisix exec deploy/apisix -- curl -s http://127.0.0.1:9080/ip -H "HOST: httpbin.local"
2 {
3   "origin": "127.0.0.1"
4 }
5 → ~ kubectl get svc -n apisix -l app.kubernetes.io/name=apisix
6 NAME                TYPE          CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
7 apisix-admin         ClusterIP     10.96.32.255    <none>           9180/TCP         13m
8 apisix-gateway       NodePort      10.96.244.106   <none>           80:31173/TCP     13m
9 → ~ curl -H "HOST: httpbin.local" \
10   `kubectl get nodes -o jsonpath="{.items[0].status.addresses[0].address}"`:31173/ip
11 {
12   "origin": "10.244.0.1"
13 }
```

总结

- Kubernetes 整体架构
- Kubernetes 核心概念
- Kubernetes 中的应用部署
- Kubernetes 中的流量路由



链接

- KIND: <https://kind.sigs.k8s.io/>
- Apache APISIX: <https://apisix.apache.org/>
- Ingress-nginx: <https://github.com/kubernetes/ingress-nginx>

