

# 应用容器化的最佳实践

张晋涛

## 个人介绍

- 张晋涛
- Apache APISIX committer
- Kubernetes ingress-nginx reviewer
- 『K8S 生态周报』发起人
- [GitHub: tao12345666333](https://github.com/tao12345666333)





01 容器化改造的必经之路

02 容器化的痛点

03 容器化的最佳实践

04 容器化扩展



# 01、容器化改造的必经之路



# 容器镜像构建

- 容器镜像是容器化交付的基本单元
- 容器镜像中包含了应用程序所需的完整依赖
- 镜像构建，交付，存储，管理等环节
- 安全性考虑

# 服务编排



- Kubernetes 已经是不二选择
- 服务间的部署架构
- 服务间的调用链

# 服务注册发现



- 基于 Kubernetes DNS 的服务注册发现机制
- 基于外部注册中心的服务注册发现机制

# 应用交付模式



- CI/CD
- GitOps
- DevSecOps



# 应用稳定性保障



- 可观测性
  - Metrics
  - Logging
  - tracing
- Debug



## 02、容器化的痛点

# 镜像构建



- 镜像体积
- 镜像构建耗时
- 密钥管理
- 供应链安全

# 镜像交付



- tag 管理
- CI/CD

# 服务发现



- 容器化业务间的服务发现
- 存量业务与容器化业务的服务发现



# 03、容器化的最佳实践



# 容器镜像是什么

- Manifest
- layers

```
→ ~ docker pull debian
Using default tag: latest
latest: Pulling from library/debian
bb7d5a84853b: Pull complete
Digest: sha256:4d6ab716de467aad58e91b1b720f0badd7478847ec7a18f66027d0f8a329a43c
Status: Downloaded newer image for debian:latest
docker.io/library/debian:latest
→ ~ mkdir -p debian-image
→ ~ docker image save -o debian-image/debian.tar debian
→ ~ ls debian-image
debian.tar
→ ~ tar -C debian-image -xf debian-image/debian.tar
→ ~ tree -I debian.tar debian-image
debian-image
├── 0d587dfbc4f4800bfe9ab08662e8396ffc37060c493f8ef24b2823fef3320df6.json
├── 3c6848d3d983bc0db9c8750311dcc9b9b5efcd71c084a8ffc1fea7ba6b3d9805
│   ├── json
│   ├── layer.tar
│   └── VERSION
├── 53b67ec39af0bd928c4a92be63ffc2c0341914b8092a4db051e3abeb34e48414
│   ├── json
│   ├── layer.tar
│   └── VERSION
├── f776cfb21b5e06bb5b4883eb15c09ab928a411476b8275293c7f96d09b90f7f9.json
├── manifest.json
└── repositories

2 directories, 10 files
```

# 容器镜像是什么



```
→ ~ cat debian-image/manifest.json | jq
[
  {
    "Config": "f776cfb21b5e06bb5b4883eb15c09ab928a411476b8275293c7f96d09b90f7f9.json",
    "RepoTags": [
      "debian:latest"
    ],
    "Layers": [
      "3c6848d3d983bc0db9c8750311dcc9b9b5efcd71c084a8ffc1fea7ba6b3d9805/layer.tar"
    ]
  },
  {
    "Config": "0d587dfbc4f4800bfe9ab08662e8396ffc37060c493f8ef24b2823fef3320df6.json",
    "RepoTags": null,
    "Layers": [
      "53b67ec39af0bd928c4a92be63ffc2c0341914b8092a4db051e3abeb34e48414/layer.tar"
    ]
  }
]
```



# 容器镜像是什么



```
→ ~ cat debian-image/f776cfb21b5e06bb5b4883eb15c09ab928a411476b8275293c7f96d09b90f7f9.json | jq -r '. |  
{rootfs: .rootfs, history: .history}'  
{  
  "rootfs": {  
    "type": "layers",  
    "diff_ids": [  
      "sha256:62a747bf1719d2d37fff5670ed40de6900a95743172de1b4434cb019b56f30b4"  
    ]  
  },  
  "history": [  
    {  
      "created": "2021-10-12T01:20:30.273959207Z",  
      "created_by": "/bin/sh -c #(nop) ADD file:aea313ae50ce6474a3df142b34d4dcba4e7e0186ea6fe55389cb2ea903b9ebbb  
in / "  
    },  
    {  
      "created": "2021-10-12T01:20:30.89167925Z",  
      "created_by": "/bin/sh -c #(nop) CMD [\"bash\"]",  
      "empty_layer": true  
    }  
  ]  
}
```



# 容器镜像的作用

- 为启动容器提供必要的文件
- 记录各层的操作
- 记录各层的配置



# 如何构建镜像

- 从镜像构建
  - 简单，直接
  - 不可追溯
- 从 Dockerfile 构建
  - 可追溯
  - 可编程

# 镜像构建



## ➤ 利用缓存

```
FROM debian

RUN apt update
RUN apt install -y openjdk-8-jdk

COPY . /app

CMD [ "java", "-jar", "/app/target/gs-spring-boot-0.1.0.jar" ]
```

## ➤ 部分拷贝

```
FROM debian

RUN apt update
RUN apt install -y openjdk-8-jdk

COPY target/gs-spring-boot-0.1.0.jar /app/

CMD [ "java", "-jar", "/app/gs-spring-boot-0.1.0.jar" ]
```

# 慎用包管理器



FROM debian

RUN apt update && apt install -y openjdk-8-jdk

COPY target/gs-spring-boot-0.1.0.jar /app/

CMD [ "java", "-jar", "/app/gs-spring-boot-0.1.0.jar" ]



FROM debian

RUN apt update && apt install -y --no-install-recommends openjdk-8-jdk

COPY target/gs-spring-boot-0.1.0.jar /app/

CMD [ "java", "-jar", "/app/gs-spring-boot-0.1.0.jar" ]



```
root@5a23eb858163:/# apt install --no-install-recommends openjdk-8-jdk | grep 'additional disk space will be used'
```

...

After this operation, 344 MB of additional disk space will be used.

^C

```
root@5a23eb858163:/# apt install openjdk-8-jdk | grep 'additional disk space will be used'
```

...

After this operation, 548 MB of additional disk space will be used.

^C

# 清理包管理器缓存



```
(MoeLove) → docker run --rm -it debian
root@cd857c3ab882:/# apt -qq update
All packages are up to date.
root@cd857c3ab882:/# du -sh /var/lib/apt/lists/
16M      /var/lib/apt/lists/
root@cd857c3ab882:/#
```

```
FROM debian
```

```
RUN apt update && apt install -y --no-install-recommends openjdk-8-jdk \
    && rm -rf /var/lib/apt/lists/*
```

```
COPY target/gs-spring-boot-0.1.0.jar /app/
```

```
CMD [ "java", "-jar", "/app/gs-spring-boot-0.1.0.jar" ]
```

# 镜像构建



- 选择合适的基础镜像
- 易用性
- 安全性

```
FROM openjdk:8-jdk-stretch

COPY target/gs-spring-boot-0.1.0.jar /app/

CMD [ "java", "-jar", "/app/gs-spring-boot-0.1.0.jar" ]
```

# 镜像构建

- 保持构建环境一致性
- 步骤分离



```
FROM maven:3.6.1-jdk-8-alpine
```

```
WORKDIR /app
```

```
COPY pom.xml /app/
```

```
COPY src /app/src
```

```
RUN mvn -e -B package
```

```
CMD [ "java", "-jar", "/app/target/gs-spring-boot-0.1.0.jar" ]
```



```
FROM maven:3.6.1-jdk-8-alpine
```

```
WORKDIR /app
```

```
COPY pom.xml /app/
```

```
RUN mvn dependency:go-offline
```

```
COPY src /app/src
```

```
RUN mvn -e -B package
```

```
CMD [ "java", "-jar", "/app/target/gs-spring-boot-0.1.0.jar" ]
```



## ➤ 多阶段构建

```
FROM maven:3.6.1-jdk-8-alpine AS builder

WORKDIR /app

COPY pom.xml /app/
RUN mvn dependency:go-offline
COPY src /app/src
RUN mvn -e -B package

FROM builder AS dev

RUN apk add --no-cache vim

FROM openjdk:8-jre-alpine

COPY --from=builder /app/target/gs-spring-boot-0.1.0.jar /

CMD [ "java", "-jar", "/gs-spring-boot-0.1.0.jar" ]
```

```
FROM maven:3.6.1-jdk-8-alpine AS builder

WORKDIR /app

COPY pom.xml /app/
RUN mvn dependency:go-offline
COPY src /app/src
RUN mvn -e -B package

FROM openjdk:8-jre-alpine

COPY --from=builder /app/target/gs-spring-boot-0.1.0.jar /

CMD [ "java", "-jar", "/gs-spring-boot-0.1.0.jar" ]
```

# 密钥管理



- 不推荐
  - 将密码硬编码写入代码中
- 一般做法
  - 通过环境变量的方式构建
- 推荐
  - Buildkit 挂载特性

```
# syntax = docker/dockerfile:experimental

COPY fetch_remote_data.sh .
RUN --mount=type=secret,id=moelove,target=/cache_builder,required ./fetch_remote_data.sh

# docker build --secret id=moelove,src=./secret -t local/spring-boot:4 .
```

# 镜像构建

## ➤ 挂载 ssh 密钥



```
(MoeLove) → d eval $(ssh-agent)
Agent pid 28184
(MoeLove) → d ssh-add ~/.ssh/id_rsa
Enter passphrase for /home/tao/.ssh/id_rsa:
Identity added: /home/tao/.ssh/id_rsa (/home/tao/.ssh/id_rsa)
(MoeLove) → d docker build --ssh=default -t local/ssh .
[+] Building 0.5s (10/10) FINISHED
=> [internal] load build definition from Dockerfile 0.1s
=> => transferring dockerfile: 96B 0.0s
=> [internal] load .dockerignore 0.1s
=> => transferring context: 2B 0.0s
=> resolve image config for docker.io/docker/dockerfile:experimental 0.0s
=> CACHED docker-image://docker.io/docker/dockerfile:experimental 0.0s
=> [internal] load metadata for docker.io/library/alpine:latest 0.0s
=> [1/4] FROM docker.io/library/alpine 0.0s
=> CACHED [2/4] RUN apk add --no-cache git openssh-client 0.0s
=> CACHED [3/4] RUN mkdir -p -m 0700 ~/.ssh && ssh-keyscan github.com >> ~/.ssh/known_hosts 0.0s
=> CACHED [4/4] RUN --mount=type=ssh,required git clone git@github.com:tao12345666333/moe.git 0.0s
=> exporting to image 0.0s
=> => exporting layers 0.0s
=> => writing image sha256:35d3ded5595a48de50054121feed13ebadf9b5e73b6cfeeba4215e1a20a20fd 0.0s
=> => naming to docker.io/local/ssh
```



```
# syntax = docker/dockerfile:experimental
FROM alpine

# 安装必要的包
RUN apk add --no-cache git openssh-client

# 创建必要的目录 .ssh 由于要使用 ssh 连接，所以需要预先使用 ssh-keyscan 先获取 public SSH host key
# 当然也可以给 .ssh/config 写配置文件来跳过验证，但容易带来安全问题，不推荐
RUN mkdir -p -m 0700 ~/.ssh && ssh-keyscan github.com >> ~/.ssh/known_hosts

# clone 私有项目仓库，并创建分支
RUN --mount=type=ssh,required git clone git@github.com:tao12345666333/moe.git \
    && cd moe \
    && git checkout -b release
```



## 04、容器化扩展

# 交付方式

- GitOps
- DevSecOps



trivy



# 策略控制



- OPA
- Gatekeeper
- Kyverno



Open Policy Agent



Kyverno



北京开课吧科技有限公司  
[www.kaikeba.com](http://www.kaikeba.com)  
400-996-0826

# Thanks