



容器镜像构建和交付的最佳实践

2022 年 Q4 版

张晋涛

个人介绍

- 张晋涛
- Apache APISIX PMC
- Kubernetes Ingress NGINX maintainer
- Microsoft MVP
- 【K8S 生态周报】维护者
- 公众号： MoeLove
- <https://github.com/tao12345666333>





目录

CONTENTS

01 • 容器镜像是什么

02 • 如何选择镜像构建工具

03 • 利用缓存提升构建效率

04 • 如何进行镜像交付



容器镜像是什么

一系列配置文件和归档文件

- 通过 skopeo 下载到本地



```
→ debian-oci skopeo copy docker://index.docker.io/library/debian
dir:./debian
Getting image source signatures
Copying blob a8ca11554fce done
Copying config c31f65dd4c done
Writing manifest to image destination
Storing signatures
→ debian-oci tree debian
debian
├── a8ca11554fce00d9177da2d76307bdc06df7faeb84529755c648ac4886192ed1
├── c31f65dd4cc97f21bd440df4b9b919a99e35a7ede602e8150f2314af1441a312
├── manifest.json
└── version

0 directories, 4 files
```



```
→ debian-oci file debian/a8ca11554fce00d9177da2d76307bdc06df7faeb84529755c648ac4886192ed1
debian/a8ca11554fce00d9177da2d76307bdc06df7faeb84529755c648ac4886192ed1: gzip compressed data, original size modulo
2^32 129330176
→ debian-oci file debian/c31f65dd4cc97f21bd440df4b9b919a99e35a7ede602e8150f2314af1441a312
debian/c31f65dd4cc97f21bd440df4b9b919a99e35a7ede602e8150f2314af1441a312: JSON text data
```

一系列配置文件和归档文件

- 通过 docker save 保存到本地

```
→ debian-image docker save -o debian.tar debian
→ debian-image mkdir -p debian
→ debian-image tar -C debian -xf debian.tar
→ debian-image tree debian
debian
├── 860b7460ac90e2e79c6e06ccae67c6bcbcaa1abbba22d37602cba3f121486e04
│   ├── json
│   ├── layer.tar
│   └── VERSION
├── c31f65dd4cc97f21bd440df4b9b919a99e35a7ede602e8150f2314af1441a312.json
├── manifest.json
└── repositories

1 directory, 6 files
```

一系列配置文件和归档文件

docker save 是解压后的内容

```
→ debian-oci jq . debian/manifest.json
{
  "schemaVersion": 2,
  "mediaType": "application/vnd.docker.distribution.manifest.v2+json",
  "config": {
    "mediaType": "application/vnd.docker.container.image.v1+json",
    "size": 1463,
    "digest":
"sha256:c31f65dd4cc97f21bd440df4b9b919a99e35a7ede602e8150f2314af1441a312"
  },
  "layers": [
    {
      "mediaType": "application/vnd.docker.image.rootfs.diff.tar.gzip",
      "size": 55038615,
      "digest":
"sha256:a8ca11554fce00d9177da2d76307bdc06df7faeb84529755c648ac4886192ed1"
    }
  ]
}
```

```
→ debian-oci docker image ls --digests debian
```

REPOSITORY	TAG	DIGEST	IMAGE ID
debian	latest	sha256:3066ef83131c678999ce82e8473e8d017345a30f5573ad3e44f62e5c9c46442b	c31f65dd4cc9 2
weeks ago	124MB		

一系列配置文件和归档文件

- 内容寻址是其核心
- 容器镜像可以由很多层进行叠加



```
→ debian-oci jq . debian/manifest.json
{
  "schemaVersion": 2,
  "mediaType": "application/vnd.docker.distribution.manifest.v2+json",
  "config": {
    "mediaType": "application/vnd.docker.container.image.v1+json",
    "size": 1463,
    "digest":
"sha256:c31f65dd4cc97f21bd440df4b9b919a99e35a7ede602e8150f2314af1441a312"
  },
  "layers": [
    {
      "mediaType": "application/vnd.docker.image.rootfs.diff.tar.gzip",
      "size": 55038615,
      "digest":
"sha256:a8ca11554fce00d9177da2d76307bdc06df7faeb84529755c648ac4886192ed1"
    }
  ]
}
```



```
→ debian-oci docker image history debian
IMAGE          CREATED          CREATED BY
SIZE           COMMENT
c31f65dd4cc9   2 weeks ago    /bin/sh -c #(nop)  CMD ["bash"]
0B
<missing>     2 weeks ago    /bin/sh -c #(nop)  ADD file:abc6873c98339a3c4...
124MB
```


如何构建容器镜像

- Dockerfile 是“构建规则”
- docker build 是最常用的命令
- ○ ○ ○

如何选择镜像构建工具

- <https://github.com/docker/docker>
- <https://github.com/containers/buildah>
- <https://github.com/GoogleContainerTools/kaniko>
- <https://github.com/ko-build/ko>
- <https://github.com/buildpacks/pack>
- <https://github.com/moby/buildkit>
- <https://github.com/genuinetools/img>(不维护)
- ...

- 当前基础设施
- 是否需要 deamon 运行
- 通用 / 专用工具
- 是否兼容 Dockerfile

- Buildkit 下一代 Docker 构建引擎
- Docker Desktop 默认启用 BuildKit
- docker 已经将默认的 docker build 切换为了 buildx(为 buildkit 所做的扩展)

- 通过 `syntax=docker/dockerfile:1` 设置语法
- 增加 `here-docs` 支持
- 书写更加简单
- 支持 shebang (`#!`) 比如 `python` 等语言

```
FROM debian

RUN apt-get update && \
    apt-get upgrade -y && \
    apt-get install -y dnsutils && \
    rm -rf /var/lib/apt/lists/*
```

```
# syntax=docker/dockerfile:1
FROM debian

RUN <<EOF
    apt-get update
    apt-get upgrade -y
    apt-get install -y dnsutils
    rm -rf /var/lib/apt/lists/*
EOF
```



利用缓存提升构建效率

Docker 有构建时缓存，已执行命令可被缓存。 COPY 命令会终止该缓存

```
+ buildkit-demo docker build -t local/debian:bad-style .
[+] Building 37.0s (10/10) FINISHED
=> [internal] load build definition from Dockerfile                                0.0s
=> => transferring dockerfile: 223B                                              0.0s
=> [internal] load .dockerignore                                                  0.0s
=> => transferring context: 2B                                                    0.0s
=> resolve image config for docker.io/docker/dockerfile:1                      2.6s
=> CACHED docker-image://docker.io/docker/dockerfile:1@sha256:9ba7531bd80fb0a858632727cf7a112fbfd19b17e94c4e84ced81e24ef1a0dbc 0.0s
=> [internal] load metadata for docker.io/library/debian:latest                 0.0s
=> CACHED [1/4] FROM docker.io/library/debian                                   0.0s
=> [2/4] RUN apt-get update                                                       22.9s
=> [3/4] RUN apt-get install -y dnsutils                                         10.2s
=> [4/4] RUN rm -rf /var/lib/apt/lists/*                                          0.3s
=> exporting to image                                                            0.7s
=> => exporting layers                                                            0.7s
=> => writing image sha256:b4931683c0f77f67118b96c6c6048fe4fea6e7aa7b7c42a79616ee6c5f8c7284 0.0s
=> => naming to docker.io/local/debian:bad-style                               0.0s
+ buildkit-demo docker build -t local/debian:bad-style .
[+] Building 1.3s (10/10) FINISHED
=> [internal] load build definition from Dockerfile                                0.0s
=> => transferring dockerfile: 223B                                              0.0s
=> [internal] load .dockerignore                                                  0.0s
=> => transferring context: 2B                                                    0.0s
=> resolve image config for docker.io/docker/dockerfile:1                      1.0s
=> CACHED docker-image://docker.io/docker/dockerfile:1@sha256:9ba7531bd80fb0a858632727cf7a112fbfd19b17e94c4e84ced81e24ef1a0dbc 0.0s
=> [internal] load metadata for docker.io/library/debian:latest                 0.0s
=> [1/4] FROM docker.io/library/debian                                           0.0s
=> CACHED [2/4] RUN apt-get update                                                0.0s
=> CACHED [3/4] RUN apt-get install -y dnsutils                                  0.0s
=> CACHED [4/4] RUN rm -rf /var/lib/apt/lists/*                                  0.0s
=> exporting to image                                                            0.0s
=> => exporting layers                                                            0.0s
=> => writing image sha256:b4931683c0f77f67118b96c6c6048fe4fea6e7aa7b7c42a79616ee6c5f8c7284 0.0s
=> => naming to docker.io/local/debian:bad-style                               0.0s
+ buildkit-demo
```



```
name: ci
on:
  push:
    branches:
      - 'main'

jobs:
  docker:
    runs-on: ubuntu-latest
    steps:
      - name: Set up Docker Buildx
        uses: docker/setup-buildx-action@v1
      - name: Login to Docker Hub
        uses: docker/login-action@v1
        with:
          username: ${ secrets.DOCKERHUB_USERNAME }
          password: ${ secrets.DOCKERHUB_TOKEN }
      - name: Build and push
        uses: docker/build-push-action@v2
        with:
          push: true
          tags: user/app:latest
```

mount=type=bind 可直接从 build context 或者 stage 进行文件操作

mount=type=cache 可进行应用的缓存

```
# syntax=docker/dockerfile:1

FROM golang:1.19-alpine AS build
WORKDIR /src
RUN --mount=type=bind,target=. \
    go build -ldflags "-s -w" -o /usr/bin/app .

FROM alpine
COPY --from=build /usr/bin/app /usr/bin/app
CMD ["app"]
```

```
name: ci
on:
  push:
    branches:
      - 'main'

jobs:
  docker:
    runs-on: ubuntu-latest
    steps:
      - name: Set up Docker Buildx
        uses: docker/setup-buildx-action@v1
      - name: Login to Docker Hub
        uses: docker/login-action@v1
        with:
          username: ${ secrets.DOCKERHUB_USERNAME }
          password: ${ secrets.DOCKERHUB_TOKEN }
      - name: Build and push
        uses: docker/build-push-action@v2
        with:
          push: true
          tags: user/app:latest
          cache-from: gha
          cache-to: gha
```

Buildkit 支持多种 cache： inline， local， registry 和 gha

```
#15 importing cache manifest from gha:2200847280167466126
#15 DONE 0.5s

#13 [internal] load build context
#13 transferring context: 1.13kB done
#13 DONE 0.0s

#16 [build 2/3] WORKDIR /src
#16 CACHED

#17 [build 3/3] RUN --mount=type=bind,target=/src go build
    -ldflags "-s -w" -o /usr/bin/app .
#17 CACHED

#18 [stage-1 2/2] COPY --from=build /usr/bin/app /usr/bin/app
#18 CACHED

#19 exporting cache
#19 preparing build cache for export done
#19 DONE 0.4s
```

通过 --link 更高效的利用
缓存
尤其是对多阶段构建
可以 --link 加 --from 配合

```
FROM alpine

COPY --link foo /

COPY --link bar /
```

Without --link

snapshot1 parent=nil			
bin/	dev/	etc/	root/
usr/	var/		

snapshot2 parent=snapshot1			
bin/	dev/	etc/	foo
root/	usr/	var/	

snapshot3 parent=snapshot2			
bar	bin/	dev/	etc/
foo	root/	usr/	var/

image			
blob(snapshot1)			
diff(snapshot1, snapshot2)			
diff(snapshot2, snapshot3)			

With --link


snapshot1 parent=nil			
bin/	dev/	etc/	root/
usr/	var/		

snapshot2 parent=nil			
foo			

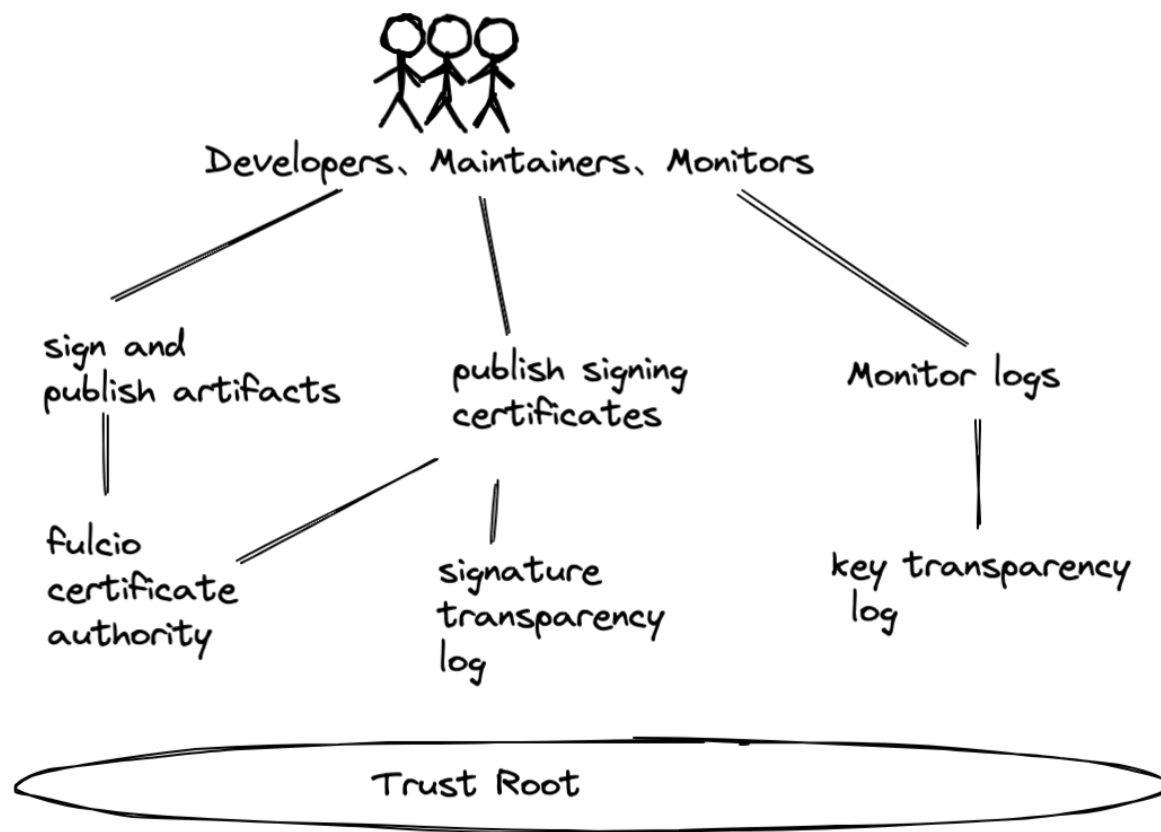
snapshot3 parent=nil			
bar			

image			
blob(snapshot1)			
blob(snapshot2)			
blob(snapshot3)			

如何进行镜像交付



使用 sigstore 和 cosign 进行签发校验





Thanks!

<https://github.com/tao12345666333>